

Chapter 6

January 12, 2021

1 Housekeeping:

```
[2]: sessionInfo()  
options(repr.plot.width=14,repr.plot.antialias='subpixel',repr.plot.res=218)  
update.packages()
```

R version 4.0.3 Patched (2020-10-12 r79333)

Platform: x86_64-apple-darwin17.0 (64-bit)

Running under: macOS Big Sur 10.16

Matrix products: default

BLAS: /Library/Frameworks/R.framework/Versions/4.0/Resources/lib/libRblas.dylib

LAPACK: /Library/Frameworks/R.framework/Versions/4.0/Resources/lib/libRlapack.dylib

locale:

[1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8

attached base packages:

[1] stats graphics grDevices utils datasets methods base

loaded via a namespace (and not attached):

[1] compiler_4.0.3 ellipsis_0.3.1 IRdisplay_0.7.0 pbdZMQ_0.3-3
[5] tools_4.0.3 htmltools_0.5.0 pillar_1.4.6 base64enc_0.1-3
[9] crayon_1.3.4 uuid_0.1-4 IRkernel_1.1.1 jsonlite_1.7.1
[13] digest_0.6.25 lifecycle_0.2.0 repr_1.1.0 rlang_0.4.8
[17] evaluate_0.14

2 Back to OJ

```
[3]: oj <- read.csv("oj.csv")
```

A log-log regression to get at elasticities:

```
[4]: basefit <- lm(log(sales) ~ log(price), data=oj)  
coef(basefit)
```

(Intercept)	10.4234223275755	log(price)	-1.60130655250769
--------------------	------------------	-------------------	-------------------

Controlling for brands in OLS:

```
[5]: brandfit <- lm(log(sales) ~ brand + log(price), data=oj)
      coef(brandfit)
```

(Intercept)	10.8288215795649	brandminute.maid	0.87017465165164	brandtropicana
	1.52994276420652	log(price)	-3.13869142792174	

Residualizing to show what partialing out things mean:

```
[6]: pricereg <- lm(log(price) ~ brand, data=oj)
      phat <- predict(pricereg, newdata=oj)
      presid <- log(oj$price)-phat
      coef( residfit <- lm( log(sales) ~ presid, data=oj) )
```

(Intercept)	9.16786399316262	presid	-3.1386914279217
--------------------	------------------	---------------	------------------

We get the same coefficient on the residual or the original variable if we control for the predicted value:

```
[7]: residfit2 <- lm( log(sales) ~ phat + log(price), data=oj)
      coef(residfit2)
```

(Intercept)	9.26268526156084	phat	3.0177588402107	log(price)	-3.13869142792193
--------------------	------------------	-------------	-----------------	-------------------	-------------------

Traditionally, and for consistency, you would also residualize the left-hand side when you residualize the right-hand side. (You can also Google the Frisch-Waugh-Lowell Theorem. Frisch also shared the first Nobel Prize in economics.)

```
[8]: salesreg <- lm(log(sales) ~ brand, data=oj)
      shat <- predict(salesreg, newdata=oj)
      sresid <- log(oj$sales) - shat
      residfit3 <- lm( sresid ~ presid - 1)
      coef(residfit3)
```

presid: -3.138691427922

3 Not controlling for abortion

Let's turn to a famous line of research, earning Levitt a Clark Medal (also Google Melissa Dell) and the Freakonomics franchise.

For this, let's load some data that is not in comma-separated plain text (csv) this time.

(skip says 'skip the first line of the file', sep="/t" says 'tab separated')

```
[9]: data <- read.table("abortion.dat", skip=1, sep="\t")
      names(data) <- c("state", "year", "pop", "y_viol", "y_prop", "y_murd",
                      "a_murd", "a_viol", "a_prop", 'prison', 'police',
```

```
'ur','inc','pov','afdc','gun','beer')
```

Some variable descriptions:

- prison: log of lagged prisoners per capita
- police: the log of lagged police per capita
- ur: the unemployment rate
- inc: per-capita income
- pov: the poverty rate
- AFDC: generosity at year t-15
- gun: dummy for concealed weapons law
- beer: beer consumption per capita

Let's drop some unusable parts of the data, and make the time and state variables

```
[10]: data <- data[!(data$state%in%c(2,9,12)),] # AK, DC, HA are strange places
data <- data[data$year>84 & data$year<98,] # incomplete data outside these years
data$pop <- log(data$pop)
t <- data$year - 85
s <- factor(data$state) ## the states are numbered alphabetically
```

Here come the designated controls:

```
[11]: controls <- data.frame(data[,c(3,10:17)])
```

Key: y is de-trended log crime rate, criminal-age-weighted prior abortion rate a is as described below. Note we also have violent and property crime versions for both.

```
[12]: y <- data$y_murd
d <- data$a_murd
```

The abortion a_ variables are weighted average of abortion rates where weights are determined by the fraction of the type of crime committed by various age groups. For example, if 60% of violent crime were committed by 18-year-olds and 40% were committed by 19 year olds in state i, the abortion rate for violent crime at time t in state i would be constructed as .6 times the abortion rate in state i at time t - 18 plus .4 times the abortion rate in state i at time t - 19. See Donohue and Levitt (2001) for further detail.

We'll just look at murder. Note that for convenience here we've made y,d,t global: they are not in controls.

Let's win a Clark Medal:

```
[13]: summary(orig <- glm(y ~ d + t + s +., data=controls) )$coef['d',]
```

```
Estimate    -0.209811946154278 Std. Error    0.0410917700855323 t value    -5.10593595061872
Pr(>|t|)      4.50592549107522e-07
```


That abortion is only one factor influencing crime in the late 1980s points out the caution required in drawing any conclusions regarding an abortion-crime link based on time series evidence alone.

As a cautionary example, let's repeat the analysis, but for cellphones rather than abortion:

```
[14]: cell <- read.csv("us_cellphone.csv")
      cellrate <- 5*cell[,2]/(1000*cell[,3]) # center on 1985 and scale by 1997-1985
```

A smart observation: there are many things that increased with similar (but non-linear) shapes over time (cellphone usage, yoga revenues, home prices, ...). Is this loading onto abortion? What if we're just fitting a quadratic trend?

```
[15]: par(mai=c(.9,.9,.1,.1))
      plot(1985:1997, tapply(d, t, mean), bty="n", xlab="year", ylab="rate", pch=21,
            bg=2)
      points(1985:1997, cellrate, bg=4, pch=21)
      legend("topleft", fill=c(2,4), legend=c("abortion", "cellphones"), bty="n")
```



```
[16]: phone <- cellrate[t+1]
      summary(tech <- glm(y ~ phone + s + t+., data=controls))$coef['phone',]
      1 - exp(-3.72e-01)
```

```
Estimate    -0.372072322217271 Std. Error    0.069322615302163 t value    -5.36725743244805
Pr(>\textbar{}t\textbar{})          1.16863650800927e-07
0.310645757475778
```

What is happening here is that murder has been increasing quadratically, and we have no other controls that do so. To be correct, you need to allow quadratic trends that could be caused by other confounding variables (e.g. technology). We also allow interaction between the controls, and interact the nation-wide phone variable with state dummies to allow for state specific tech adoption.

```
[17]: t <- factor(t)
      interact <- glm(y ~ d + t + phone*s + .^2, data=controls)
```

```
summary(interact)$coef["d",]
```

Estimate	0.279710654345324	Std. Error	0.180731113086339	t value	1.54766188050698
Pr(>\textbar{}t\textbar{})			0.122375348876949		

The abortion effect has switched direction (and is insignificant)!

But we have very few observations relative to number of parameters:

```
[18]: dim(model.matrix(formula(interact), data=controls))
```

```
1. 624 2. 154
```

So we need a way to select only important controls. Try using a lasso (after refactoring state to have NA reference level):

```
[19]: library(gamlr)
sna <- factor(s, levels=c(NA,levels(s)), exclude=NULL)
x = sparse.model.matrix( ~ t + phone*sna + .^2, data=controls)[,-1]
dim(x)

## naive lasso regression
naive <- cv.gamlr(cbind(d,x),y)
coef(naive)["d",] # effect is CV selected <0
```

Loading required package: Matrix

```
1. 624 2. 154
```

```
-0.118963279334163
```

To explicitly include confounder that confound, we first need a model to predict treatment (this is step i of Algorithm 14):


```
[20]: treat <- cv.gamlr(x,d, lmr=1e-4)
```

Save the predicted treatment. type="response" is redundant here (Gaussian model), but you'd want it if d was binary

```
[21]: dhat <- drop( predict(treat, x, select="min") )
```

Worrying sign: not much signal in d not predicted by dhat (that means we have little to resemble an experiment here...)

```
[22]: par(mai=c(.9,.9,.1,.1))
plot(dhat,d,bty="n",pch=21,bg=8, cex=.8, yaxt="n")
axis(2, at=c(0,1,2,3))
```



Check an In Sample R^2 , which is what governs how much independent signal you have for estimating (which you see in the very similar OLS regressions below)

```
[23]: cor(drop(dhat),d)^2
```

```
0.989617332146834
```

```
[24]: summary( glm( y ~ I(d-dhat) ) )
summary( glm( y ~ d + dhat) )
```

Call:

```
glm(formula = y ~ I(d - dhat))
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-3.5530	-0.4548	0.0899	0.5427	1.2888

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-2.87706	0.02626	-109.572	<2e-16 ***
I(d - dhat)	0.32530	0.55353	0.588	0.557

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 0.4302084)

Null deviance: 267.74 on 623 degrees of freedom
Residual deviance: 267.59 on 622 degrees of freedom

AIC: 1248.5

Number of Fisher Scoring iterations: 2

Call:

```
glm(formula = y ~ d + dhat)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-3.5476	-0.4438	0.0784	0.5415	1.3061

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-2.9080	0.0321	-90.601	<2e-16 ***
d	0.3164	0.5528	0.572	0.567
dhat	-0.2218	0.5562	-0.399	0.690

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 0.4289767)

Null deviance: 267.74 on 623 degrees of freedom
Residual deviance: 266.39 on 621 degrees of freedom
AIC: 1247.7

Number of Fisher Scoring iterations: 2

Properly following Algorithm 14, we can re-run the lasso, with this prediction (the second column) included unpenalized.

```
[25]: causal <- gamlr(cbind(d,dhat,x),y,free=2,lmr=1e-3)
coef(causal)["d",] # AICc says abortion has no causal effect.
```

0

Jumping ahead to other methods in the chapter, let's try orthoML on the abortion data. (with only 5 folds, the results are a bit unstable, but at least it does not matter for the significance of the effect)

```
[26]: source("orthoML.R")
dreg <- function(x,d){ cv.gamlr(x, d, lmr=1e-5) }

yreg <- function(x,y){ cv.gamlr(x, y, lmr=1e-5) }

resids <- orthoPLTE( x=x, d=d, y=y,
                    dreg=dreg, yreg=yreg, nfold=5)
```

Loading required package: AER

Loading required package: car

Loading required package: carData

Loading required package: lmttest

Loading required package: zoo

Attaching package: 'zoo'

The following objects are masked from 'package:base':

as.Date, as.Date.numeric

Loading required package: sandwich

Loading required package: survival

fold: 1 2 3 4 5

gamma (se) = 0.0180822 (0.147449)

Traditional two-sided p-value for the key treatment effect:

```
[27]: 2*(1-pnorm(0.120546/0.164335))
```

0.463230238315691

4 Sample-splitting and orthoML in the hockey data

```
[28]: library(gamlr) # loads Matrix as well
data(hockey) # load the data
x <- cbind(config,team,player) # Combine the covariates all together
y <- goal$homegoal # build 'y': home vs away, binary response
fold <- sample.int(2,nrow(x),replace=TRUE)
nhlprereg <- gamlr(x[fold==1,], y[fold==1],
  free=1:(ncol(config)+ncol(team)),
  family="binomial", standardize=FALSE)
```

We can just take the selected variables (without the intercept, thus the -1), and fit an unpenalized model using them in the other half of the data.

```
[29]: selected <- which(coef(nhlprereg)[-1,] != 0)
xnotzero <- as.data.frame(as.matrix(x[,selected]))
```



```
nhlmle <- glm( y ~ ., data=xnotzero,
              subset=which(fold==2), family=binomial )
```

We can then interpret the coefficients and the (robust!) standard errors as causally.

```
[30]: x[1,x[1,]!=0]
p <- predict(nhlmle, xnotzero[1,,drop=FALSE], type="response", se.fit=TRUE)
p$fit
p$se.fit
library(AER)
Vnhl <- vcovHC(nhlmle)
```

```
DAL.20022003 -1 EDM.20022003 1 ERIC\_BREWER 1 JASON\_CHIMERA 1 ROB\_DIMAIO
-1 DERIAN\_HATCHER      -1 NIKO\_KAPANEN      -1 JERE\_LEHTINEN      -1
JUSSI\_MARKKANEN 1 JANNE\_NIINIMAA 1 RYAN\_SMYTH 1 BRIAN\_SWANSON 1
MARTY\_TURCO           -1 SERGEI\_ZUBOV           -1
```

Warning message in predict.lm(object, newdata, se.fit, scale = residual.scale,
type = if (type == :

“prediction from a rank-deficient fit may be misleading”

1: 0.505982487615667

1: 0.068946009155856

Why did we get different numbers from the book? Could the random fold matter this much? Not really. Note that we use the original hockey specification here, when we don’t penalize the configuration and team dummies. The book uses the slightly simpler example. (It is a bit surprising it matters this much.)

Let’s pick a player as an example:

```
[31]: WHO <- "SIDNEY_CROSBY"
Bhat <- coef(nhlmle)
exp(Bhat["SIDNEY_CROSBY"])
who <- grep(WHO,colnames(x))
```

SIDNEY_CROSBY: 1.60205026274221

We can repeat the propensity adjustment exercise:

```
[ ]: treatreg <- gamlr(x[,-who], x[,who],
                      free=1:(ncol(config)+ncol(team)),
                      standardize=FALSE, lmr=1e-6)
plot(treatreg) # basically we're getting OLS here; signal is strong
```

```
[33]: dhat <- predict(treatreg, x[,-who], type="response")
causereg <- gamlr(cbind(config,team,dhat,player), y,
                  free=1:(ncol(config)+ncol(team)+1),
                  family="binomial", standardize=FALSE)
coef(causereg)[WHO,]
```

1.11799427279508

Or just stick to ortogonal ML:

```
[34]: source("orthoML.R")
dreg <- function(x,d){
  gamlr(x, d, standardize=FALSE, lmr=1e-5) }

yreg <- function(x,d){
  gamlr(x, d, family="binomial", standardize=FALSE, lmr=1e-5) }

resids <- orthoPLTE( x=x[, -who], d=x[, who], y=y,
                    dreg=dreg, yreg=yreg, nfold=5)
```

```
fold: 1 2 3 4 5
gamma (se) = 0.255991 (0.0206741)
```

5 Heterogenous Treatment Effects

Back to the Oregon data:

```
[35]: library(foreign)

descr <- read.dta("OHIE_Public_Use_Files/OHIE_Data/oregonhie_descriptive_vars.
  →dta")
prgm <- read.dta("OHIE_Public_Use_Files/OHIE_Data/oregonhie_stateprograms_vars.
  →dta")
s12 <- read.dta("OHIE_Public_Use_Files/OHIE_Data/oregonhie_survey12m_vars.dta")

# nicely organized, one row per person
all(s12$person_id == descr$person_id)
all(s12$person_id == prgm$person_id)

P <- descr[,c("person_id", "household_id", "numhh_list")]
P$medicaid <- as.numeric(prgm[, "ohp_all_ever_firstn_30sep2009"] == "Enrolled")
P$selected <- as.numeric(descr[, "treatment"] == "Selected")
levels(P$numhh_list) <- c("1", "2", "3+")

# 12 month is the survey that really matters
# need to control for household size interacted with survey return time
Y <- s12[,c("weight_12m",
  "doc_any_12m", "doc_num_mod_12m",
  "er_any_12m", "er_num_mod_12m",
  "hosp_any_12m", "hosp_num_mod_12m")]
Y$doc_any_12m <- as.numeric(Y$doc_any_12m == "Yes")
Y$er_any_12m <- as.numeric(Y$er_any_12m == "Yes")
Y$hosp_any_12m <- as.numeric(Y$hosp_any_12m == "Yes")
```

```

# smk_ever_12m - num19_12m are sources of heterogeneity, plus descr
X <- s12[,121:147]
X$dt_returned <- factor(format(s12$dt_returned_12m, "%Y-%m"))

insurv <- which(s12$sample_12m_resp == "12m mail survey responder")
X <- X[insurv,]
Y <- Y[insurv,]
P <- P[insurv,]

sapply(Y,function(y) sum(is.na(y)))
nomiss <- which( !apply(Y,1, function(y) any(is.na(y))) )
X <- X[nomiss,]
Y <- Y[nomiss,]
P <- P[nomiss,]

# pull out the weights and attach doc_any to P
weights <- Y[,1]
Y <- Y[,-1]

# replace some ridiculous values in survey and drop num19
X$hhsz_12m[X$hhsz_12m>10] <- 10
X$num19_12m <- NULL

# organize to make it pretty for text
P$doc_any_12m <- Y$doc_any_12m # you can explore other responses if you want
P <- P[,c(1,2,6,5,4,3)]
names(P)[6] <- "numhh"

```

TRUE

TRUE

```

weight\_12m    0 doc\_any\_12m    249 doc\_num\_mod\_12m    300 er\_any\_12m    227
er\_num\_mod\_12m    260 hosp\_any\_12m    168 hosp\_num\_mod\_12m    202

```

5.1 Dealing with missingness

As often, make NA the reference level for all categories

```

[36]: source("naref.R")
X <- naref(X)

```

Now take the numeric variables, flag them when they are missing, and impute zero for sparse variables and means for dense ones.

```

[37]: xnum <- X[,sapply(X,class)%in%c("numeric","integer")]
xnum[66:70,]
colSums(is.na(xnum))
# flag missing

```

```

xnumna <- apply(is.na(xnum), 2, as.numeric)
xnumna[66:70,]
# impute the missing values
mzimpute <- function(v){
  if(mean(v==0,na.rm=TRUE) > 0.5) impt <- 0
  else impt <- mean(v, na.rm=TRUE)
  v[is.na(v)] <- impt
  return(v) }
xnum <- apply(xnum, 2, mzimpute)
xnum[66:70,]

# replace/add the variables in new data frame
for(v in colnames(xnum)){
  X[,v] <- xnum[,v]
  X[,paste(v,"NA", sep=".")] <- xnumna[,v] }
X[144:147,]

```

		smk_avg_mod_12m <int>	birthyear_12m <int>	hhinc_pctfpl_12m <dbl>	hhsz_12m <dbl>
A data.frame: 5 × 4	66	0	1974	NA	NA
	67	15	1963	150.04617	1
	68	NA	1962	150.04617	1
	69	20	1964	61.44183	3
	70	10	NA	14.71825	10

smk_avg_mod_12m 3356 birthyear_12m 518 hhinc_pctfpl_12m 2253 hhsz_12m 1175

		smk_avg_mod_12m	birthyear_12m	hhinc_pctfpl_12m	hhsz_12m
A matrix: 5 × 4 of type dbl	0	0	1	1	
	0	0	0	0	
	1	0	0	0	
	0	0	0	0	
	0	1	0	0	

		smk_avg_mod_12m	birthyear_12m	hhinc_pctfpl_12m	hhsz_12m
A matrix: 5 × 4 of type dbl	0	1974.000	77.20707	2.987188	
	15	1963.000	150.04617	1.000000	
	0	1962.000	150.04617	1.000000	
	20	1964.000	61.44183	3.000000	
	10	1965.777	14.71825	10.000000	

		smk_ever_12m <fct>	smk_curr_12m <fct>	smk_avg_mod_12m <dbl>	smk_quit_12m <fct>	female_12 <fct>
A data.frame: 4 × 31	144	No	not at all	0	NA	Male
	145	Yes	every day	14	No	Male
	146	Yes	every day	6	Yes	Female
	147	Yes	not at all	0	NA	Female

Now we just build the usual design matrix.

```
[38]: xhte <- sparse.model.matrix(~., data=cbind(numhh=P$numhh, X))[, -1]
xhte[1:2, 1:4]
dim(xhte)
```

```
2 x 4 sparse Matrix of class "dgCMatrix"
  numhh2 numhh3+ smk_ever_12mNo smk_ever_12mYes
1      .      .              .              1
2      .      .              1              .

1. 23107 2. 91
```

And interactions for an even larger design matrix:

```
[39]: dxhte <- P$selected*xhte
colnames(dxhte) <- paste("d", colnames(xhte), sep=".")
htedesign <- cbind(xhte, d=P$selected, dxhte)
```

And now we can fit the interacted model, only we need to include the baseline treatment and its numhh controls without penalty.

```
[40]: htefit <- gamlr(x=htedesign, y=P$doc_any_12m, free=c("numhh2", "numhh3+", "d"))
gam <- coef(htefit)[- (1:(ncol(xhte)+1)), ]
round(sort(gam)[1:6], 4)
round(sort(gam, decreasing=TRUE)[1:6], 4)
```

```
[1] 1 2 92

d.race\_asian\_12mYes      -0.0446 d.employ\_hrs\_12mwork 20-29 hrs/week      -0.0433
d.hhinc\_cat\_12m\$32501-\$35000  -0.0293 d.hhinc\_cat\_12m\$27501-\$30000  -0.0232
d.hhinc\_cat\_12m\$15001-\$17500  -0.0195 d.race\_hisp\_12mYes      -0.0173

d      0.0927 d.race\_pacific\_12mYes      0.0404 d.hhinc\_cat\_12m\$2501-\$5000      0.0221
d.hhinc\_cat\_12m\$5001-\$7500      0.0137 d.live\_other\_12mYes      0.0116
d.race\_black\_12mYes              0.0067
```

There is one extra complication here: as the interacting variables don't average to a zero shift of the treatment effect, the ATE is not just the coefficient on the *d* variable but also the weighted sum of the heterogeneous treatment effects. They can be collected by hand, or with a shorthand that certainly takes some time to learn.

```
[41]: gam["d"] + colMeans(xhte)%*%gam[grep("d.", names(gam))]
```

A matrix: 1 × 1 of type dbl 0.05616546

6 Hold my beer

Load, explore and downsample the (cleaned) Dominick's beer data:

```
[42]: load("dominicks-beer.rda")
```

```

# how many upcs?
length( upctab <- table(wber$UPC) )

# check data types
sapply(wber, class)

# create priceperoz
wber$lp <- log(12*wber$PRICE/upc[wber$UPC,"OZ"])

# smallbeer
set.seed(888)
ss <- sample.int(nrow(wber),5e3)

```

287

STORE 'character' UPC 'character' WEEK 'factor' PRICE 'numeric' MOVE 'integer'

See how you can get a small elasticity:

```
[43]: coef( margfit <- lm(log(MOVE) ~ lp, data=wber[ss,]) )
```

(Intercept)	1.04608661196579 lp	-0.658260564487543
-------------	---------------------	--------------------

Now, we better have some controls. Say, dummy indicators for beer type, week, and store:

```
[44]: wber$s <- factor(wber$STORE)
wber$u <- factor(wber$UPC)
wber$w <- factor(wber$WEEK)
xs <- sparse.model.matrix( ~ s-1, data=wber)
xu <- sparse.model.matrix( ~ u-1, data=wber)
xw <- sparse.model.matrix( ~ w-1, data=wber)
```

We'll also collect some indicators for words in the description, and collect all controls:

```
[45]: #install.packages("tm")
library(tm)
descr <- Corpus(VectorSource(as.character(upc$DESCRIP)))
descr <- DocumentTermMatrix(descr)
descr <- sparseMatrix(i=descr$i,j=descr$j,x=as.numeric(descr$v>0), # convert
  ↳from stm to Matrix format
  dims=dim(descr),dimnames=list(rownames(upc),colnames(descr)))

descr[1:5,1:6]
descr[287,descr[287,]!=0]

controls <- cbind(xs, xu, xw, descr[wber$UPC,])
dim(controls)
```

Loading required package: NLP

```

5 x 6 sparse Matrix of class "dgCMatrix"
      32oz beer budweiser n.r.b 6pk busch
1820000008  1    1          1    1    .    .
1820000016  .    1          1    .    1    .
1820000051  .    1          .    .    1    1
1820000106  .    1          1    .    1    .
1820000117  1    1          1    .    .    .

6pk      1 red      1 ale      1 honey      1 oregon      1
1.1600572 2.851

```

Here's the naive lasso. (Why don't we standardize here?)

```

[46]: naivfit <- gamlr(x=cbind(lp=wber$lp,controls)[ss,], y=log(wber$MOVE)[ss],
  ↪free=1, standardize=FALSE)
print( coef(naivfit)[1:2,] )

```

```

intercept      lp
0.128270 -1.980717

```

Let's do the double/orthogonal ML:

```

[47]: source("orthoML.R")
dreg <- function(x,d){
  gamlr(x, d, standardize=FALSE, lmr=1e-5) }

yreg <- function(x,y){
  gamlr(x, y, standardize=FALSE, lmr=1e-5) }

resids <- orthoPLTE( x=controls[ss,], d=wber$lp[ss], y=log(wber$MOVE)[ss],
  ↪dreg=dreg, yreg=yreg, nfold=5)

```

```

fold: 1  2  3  4  5
gamma (se) = -3.60396 (0.121278)

```

Now here's a new trick: Actually, we can fit a big model after all, without penalty, but we would still need these fancy commands. Why? Because the standard commands work with matrices, which are horrendously large here if you want to include all the zeros everywhere. So let's you a trusty command that works with sparse matrices, only without penalties this time.

```

[48]: fullfit <- gamlr(x=cbind(lp=wber$lp,controls), y=log(wber$MOVE), lambda.start=0)
print( coef(fullfit)["lp",] )

```

```

[1] -3.567488

```

Heterogenous treatment effects are great for beer, but complicated to build.

```

[49]: # interact items and text with price
lpxu <- xu*wber$lp
colnames(lpxu) <- paste("lp",colnames(lpxu),sep="")

```

```
# create our interaction matrix
xhte <- cbind(BASELINE=1,descr[wber$UPC,])
d <- xhte*wber$lp
colnames(d) <- paste("lp",colnames(d),sep=":")


eachbeer <- xhte[match(rownames(upc),wber$UPC),]
rownames(eachbeer) <- rownames(upc)
```

Use these interactions in the fanciest model first:

```
[50]: dmlhte <- gamlr(x=xhte[ss,]*resids$dtl, y=resids$ytl, free=1,
  →standardize=FALSE)
coef(dmlhte)[1:2]
range( gamdml <- drop(eachbeer%*%coef(dmlhte)[-1,]) )
hist(gamdml, main="", xlab="elasticity", col="lightblue", freq=FALSE)
```

1. -0.00172525179814379 2. -2.83651319832095

1. -8.43628254554082 2. -0.521001080730426



```
[51]: B <- coef(dmlhte)[- (1:2),]
B <- B[B!=0]
head(sort(round(B,2)))
head(sort(round(B,2), decreasing=TRUE))
```


bonus	-4.22	lager	-3.04	drft	-2.24	draft	-1.44	longneck	-1.25	30pk	-1.19
sharp's	2.83	state	1.62	n.a.	1.55	schlitz	1.41	dry	1.23	baderbrau	1.2

```
[52]: upc[names(sort(gamdml)[1:3]),]
```


		DESCRIP <chr>	OZ <dbl>
A data.frame: 3 × 2	7336097301	OLD STYLE LT BONUS 6 24pk	288
	7336097305	OLD STYLE LT BONUS 6 24pk	288
	3410015306	MILLER GENUINE DRFT 24pk	288


The other models performs much less sensibly with all these interactions, even the MLE that looked OK before.

```
[53]: d <- xhte*wber$lp
colnames(d) <- paste("lp", colnames(d), sep=":")
naivehte <- gamlr(x=cbind(d, controls)[ss,],
                  y=log(wber$MOVE)[ss],
                  free=1, standardize=FALSE)
gammaive <- drop(eachbeer%%coef(naivehte)[2:(ncol(d)+1),])
hist(gammaive, main="", xlab="elasticity", col="lightyellow", freq=FALSE)
```




```
[54]: mlehte <- gamlr(x=cbind(d, controls)[ss,],
                    y=log(wber$MOVE)[ss], lambda.start=0)
gammle <- drop(eachbeer%%coef(mlehte)[2:(ncol(d)+1),])
hist(gammle, main="", xlab="elasticity", breaks=200, col="pink", xlim=c(-60,25),
     ↪ freq=FALSE)
sort(gammle)[1:4]
```

```
73692011121 -918.809201244423 7336043668 -23.695480005504 8797501350 -11.1158949438684
7336097301 -10.3412562421114
```



If you use the entire data, even MLE could work (under ignorability):

```
[64]: fullhte <- gamlr(x=cbind(d,controls), y=log(wber$MOVE), lambda.start=0)
      gamfull <- drop(eachbeer%%coef(fullhte)[2:(ncol(d)+1),])
      hist(gamfull, main="", xlab="elasticity", ,
           col="darkgrey", freq=FALSE)
```




But this is not a great sign:

```
[56]: upc[which(gamfull>0),]
```

		DESCRIP <chr>	OZ <dbl>
A data.frame: 5 × 2	3410000304	MILLER SHARP'S N/A C 6pk	72
	3410010505	MILLER SHARP'S N/A L 6pk	72
	5230000035	OLD MILWAUKEE N.A. C 6pk	72
	5230001300	OLD MILWAUKEE N.A. B 6pk	72
	5230001301	OLD MILWAUKEE N.A. B 6pk	72

Compare the elasticity estimates:

```
[57]: ylim <- c(-8,1)
par(mai=c(.7,.7,.1,.1), mfrow=c(1,3))
plot(gamfull, gammle, pch=21, bg="pink", xlab="fulldata MLE", ylab="subsample_
→MLE", bty="n", ylim=ylim)
text(x=-6,y=1, sprintf("R2 = %.02f",summary(lm(gamfull~gammle))$r.squared))
plot(gamfull, gamnaive, pch=21, bg="lightyellow", xlab="fulldata MLE",
→ylab="subsample Naive ML", bty="n", ylim=ylim)
text(x=-6,y=1, sprintf("R2 = %.02f",summary(lm(gamfull~gamnaive))$r.squared))
plot(gamfull, gamdml, pch=21, bg="lightblue", xlab="fulldata MLE",
→ylab="subsample Orthogonal ML", bty="n", ylim=ylim)
text(x=-6,y=1, sprintf("R2 = %.02f",summary(lm(gamfull~gamdml))$r.squared))
```



7 Synthetic controls

You are welcome to tune into the Chamberlain Seminar, by the way.

Just work with the OG data for this:

```
[58]: #install.packages(c("Synth", "tidyr"))
library(Synth)
library(tidyr)
data(basque)
```

```
##
```

```
## Synth Package: Implements Synthetic Control Methods.
```

```
## See http://www.mit.edu/~jhainm/software.htm for additional information.
```

Attaching package: ‘tidyr’

The following object is masked _by_ ‘.GlobalEnv’:

who

The following objects are masked from ‘package:Matrix’:

expand, pack, unpack

Collect outcomes for treated region pre and post, and for all control.

```
[59]: y <- basque[,1:4] %>% spread(year, gdpcap)
rownames(y) <- y$regionname
y <- y[c(17,2:16,18), -(1:2)]
y <- y[,1:35]
```


We can define a function that runs an ML prediction routine to predict a control. This hard codes the treated year here.

```
[60]: library(gamlr)
synthc <- function(j, tyear=1968, ...){
  y0t <- t(y[,1:(tyear-1954)])
  fit <- gamlr( y0t[,-j], y0t[,j], ...)
  plot(fit)
  y0hat <- predict(fit, t(y[-j,]))[,1]
  return(list(w=coef(fit)[,1], y0hat=y0hat ) )
}
```

And run it.

```
[61]: sc <- synthc(1, lmr=1e-4)
      sc$w
```

```
intercept 0.832492340851006 Andalusia 0 Aragon 0 Principado De Asturias 0 Baleares (Islas)
-0.143662241431167 Canarias 0 Cantabria 0 Castilla Y Leon 0.515992402831998 Castilla-La
Mancha 0 Cataluna 0 Comunidad Valenciana -0.358177982636518 Extremadura 0 Galicia 0
Madrid (Comunidad De) 0.426397880130771 Murcia (Region de) 0.892844824789896 Navarra
(Comunidad Foral De) 0 Rioja (La) 0
```



We can do a permutation test for untreated regions, and this is typically a task that can be parallelized.

```
[62]: library(parallel)
      cl <- makeCluster(detectCores())
      clusterExport(cl, c("y", "gamlr", "synthc"))

      gety0 <- function(j){ synthc(j, lmr=1e-4)$y0hat }
      Ysynth <- parSapply(cl, 1:nrow(y), gety0)
      diff <- Ysynth - t(y)
```

Aaaaand plots!

```
[63]: year <- as.numeric(colnames(y))
      plot(year, sc$y0hat, type="l", ylab="gdp per capita",
            col=rgb(.1,.5,1,0.8), ylim=range(c(y[1,],sc$y0hat)), bty="n", lwd=2)
      abline(v=1968, col=8, lty=2)
      lines(year, y[1,], col=rgb(1,.5,0,.8), lwd=2)
      legend("topleft", bty="n", legend=c("observed basque","synthetic basque"),
            lwd=2, col=c(col=rgb(1,.5,0,.8),rgb(.1,.5,1,0.8)) )
```

```

matplot(year, diff, type="l", lwd=1.5,
        xlab="year", ylab="synthetic - observed",
        col=8, lty=1, bty="n")
lines(year, diff[,1], lwd=1.5, col="red")
lines(year, diff[,14], lwd=1.5, lty=2, col=1)
legend("topleft", bty="n",
      legend=c("basque", "placebo", "(madrid)"),
      lty=c(1,1,2), lwd=2, col=c(2,8,1))

```

